

Homework_3

Akeem Ajede

10/17/2019

Question 10

(a)

```
## Numerical Summary
library(ISLR)
dim(Weekly)

## [1] 1089     9

str(Weekly)

## 'data.frame': 1089 obs. of 9 variables:
## $ Year      : num  1990 1990 1990 1990 1990 ...
## $ Lag1      : num  0.816 -0.27 -2.576 3.514 0.712 ...
## $ Lag2      : num  1.572 0.816 -0.27 -2.576 3.514 ...
## $ Lag3      : num  -3.936 1.572 0.816 -0.27 -2.576 ...
## $ Lag4      : num  -0.229 -3.936 1.572 0.816 -0.27 ...
## $ Lag5      : num  -3.484 -0.229 -3.936 1.572 0.816 ...
## $ Volume    : num  0.155 0.149 0.16 0.162 0.154 ...
## $ Today     : num  -0.27 -2.576 3.514 0.712 1.178 ...
## $ Direction: Factor w/ 2 levels "Down","Up": 1 1 2 2 2 1 2 2 2 1 ...

summary(Weekly)

##      Year          Lag1          Lag2          Lag3      
## Min. :1990   Min. :-18.1950   Min. :-18.1950   Min. :-18.1950  
## 1st Qu.:1995  1st Qu.: -1.1540  1st Qu.: -1.1540  1st Qu.: -1.1580 
## Median :2000  Median : 0.2410  Median : 0.2410  Median : 0.2410  
## Mean   :2000  Mean   : 0.1506  Mean   : 0.1511  Mean   : 0.1472  
## 3rd Qu.:2005  3rd Qu.:  1.4050  3rd Qu.:  1.4090  3rd Qu.:  1.4090 
## Max.  :2010  Max.  : 12.0260  Max.  : 12.0260  Max.  : 12.0260  
##      Lag4          Lag5          Volume        
## Min. :-18.1950  Min. :-18.1950  Min. :0.08747  
## 1st Qu.: -1.1580 1st Qu.: -1.1660  1st Qu.:0.33202 
## Median : 0.2380  Median : 0.2340  Median :1.00268  
## Mean   : 0.1458  Mean   : 0.1399  Mean   :1.57462  
## 3rd Qu.:  1.4090  3rd Qu.:  1.4050  3rd Qu.:2.05373 
## Max.  : 12.0260  Max.  : 12.0260  Max.  :9.32821  
##      Today        Direction    
## Min. :-18.1950  Down:484    
## 1st Qu.: -1.1540 Up  :605    
## Median : 0.2410
```

```

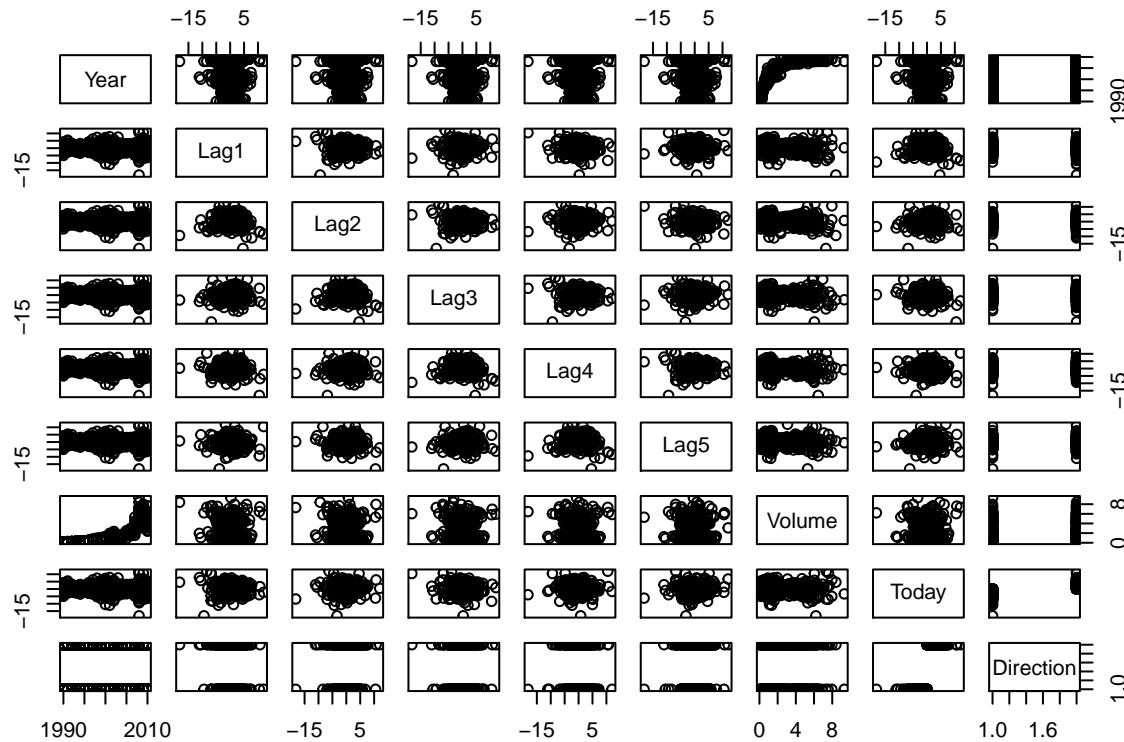
##  Mean   :  0.1499
##  3rd Qu.:  1.4050
##  Max.   : 12.0260

```

```

## Graphical Summary
pairs(Weekly)

```



```

cor(Weekly[,1:8]) #There is a strong positive relationship between "Year" and "Volume". lag2 and Lag5 s

```

```

##          Year      Lag1      Lag2      Lag3      Lag4
## Year  1.00000000 -0.032289274 -0.03339001 -0.03000649 -0.031127923
## Lag1  -0.03228927  1.000000000 -0.07485305  0.05863568 -0.071273876
## Lag2  -0.03339001 -0.074853051  1.00000000 -0.07572091  0.058381535
## Lag3  -0.03000649  0.058635682 -0.07572091  1.00000000 -0.075395865
## Lag4  -0.03112792 -0.071273876  0.05838153 -0.07539587  1.000000000
## Lag5  -0.03051910 -0.008183096 -0.07249948  0.06065717 -0.075675027
## Volume 0.84194162 -0.064951313 -0.08551314 -0.06928771 -0.061074617
## Today -0.03245989 -0.075031842  0.05916672 -0.07124364 -0.007825873
##          Lag5      Volume     Today
## Year  -0.030519101  0.84194162 -0.032459894
## Lag1  -0.008183096 -0.06495131 -0.075031842
## Lag2  -0.072499482 -0.08551314  0.059166717
## Lag3  0.060657175 -0.06928771 -0.071243639
## Lag4  -0.075675027 -0.06107462 -0.007825873
## Lag5  1.000000000 -0.05851741  0.011012698

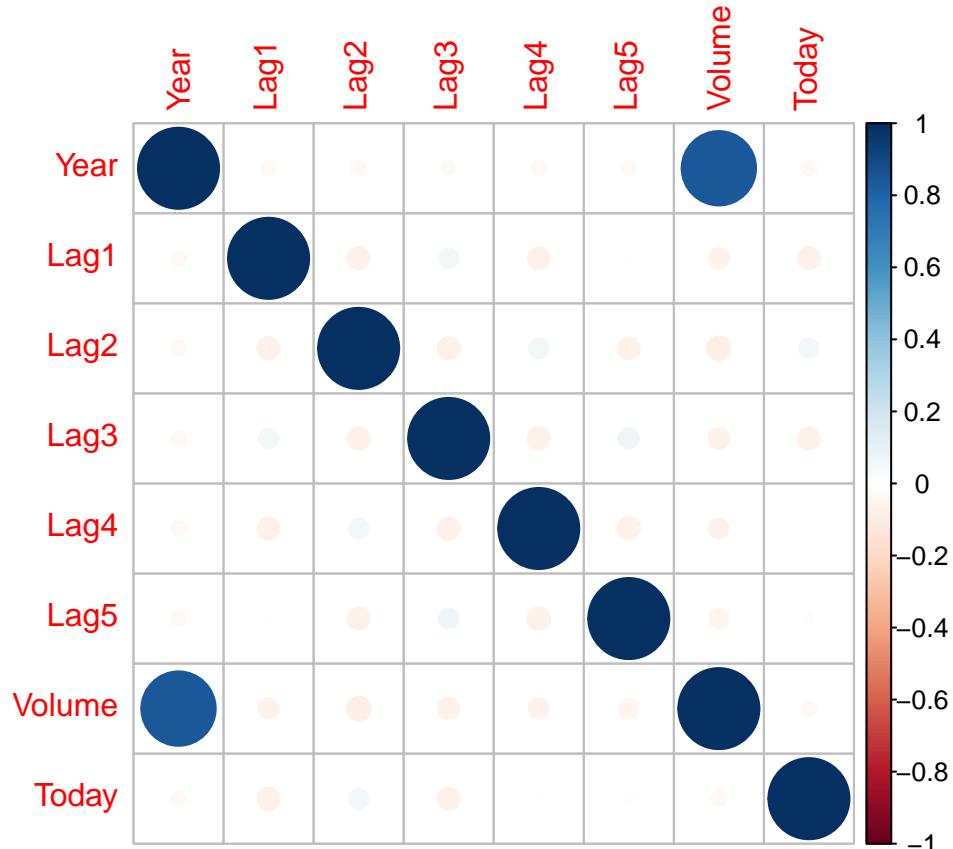
```

```
## Volume -0.058517414 1.000000000 -0.033077783  
## Today 0.011012698 -0.03307778 1.000000000
```

```
library(corrplot)
```

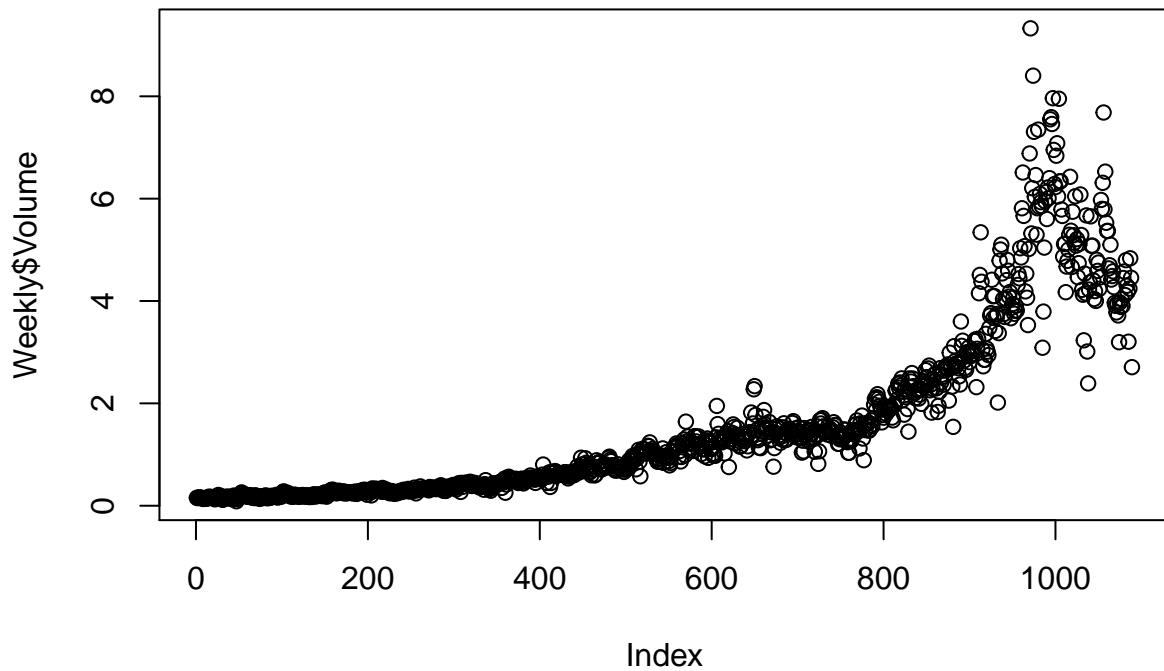
```
## corrplot 0.84 loaded
```

```
corrplot(cor(Weekly[,1:8]))
```

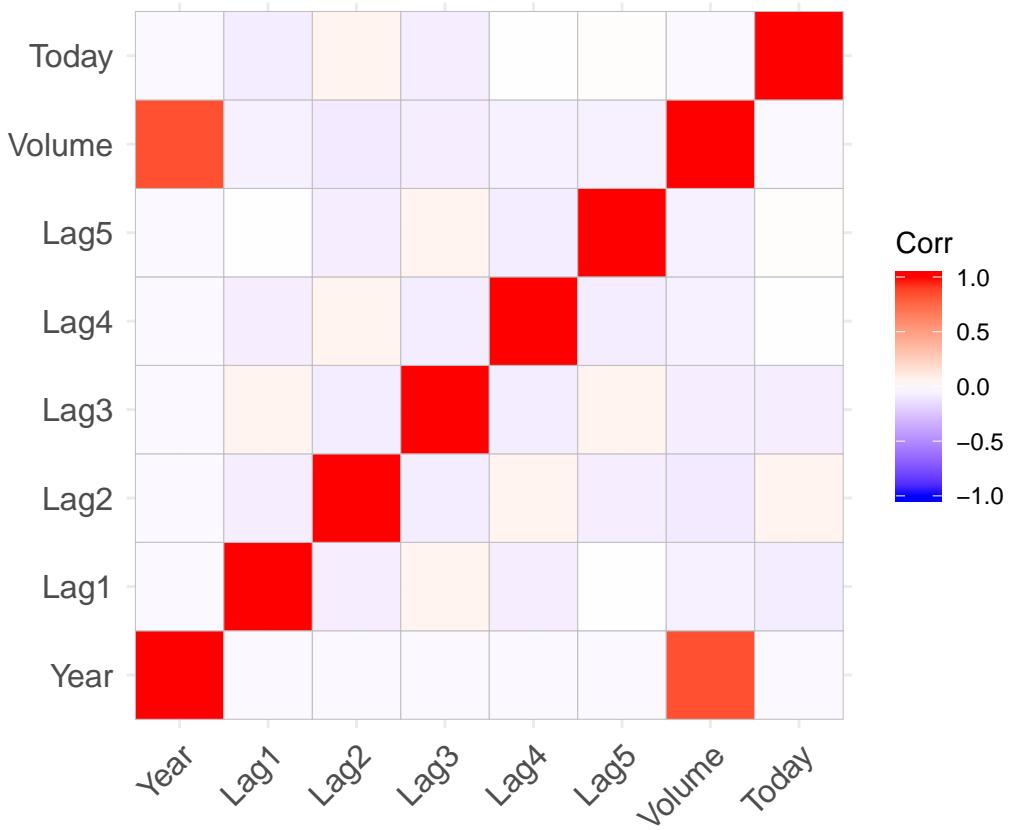


```
plot(Weekly$Volume) #Corroborates the previous deduction about the "Year" and "Volume" correlation.  
library(ggcorrplot)
```

```
## Loading required package: ggplot2
```



```
ggcorrplot(cor(Weekly[,1:8])) #Another approach to the dataset visualization
```



(b)

```
attach(Weekly)
glm.fit = glm(Direction~Lag1+Lag2+Lag3+Lag4+Lag5+Volume, data = Weekly,
              family = binomial)
summary(glm.fit)
```

```
##
## Call:
## glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 +
##       Volume, family = binomial, data = Weekly)
##
## Deviance Residuals:
##      Min        1Q     Median        3Q       Max
## -1.6949   -1.2565    0.9913    1.0849    1.4579
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept) 0.26686   0.08593   3.106   0.0019 **
## Lag1        -0.04127   0.02641  -1.563   0.1181
## Lag2         0.05844   0.02686   2.175   0.0296 *
## Lag3        -0.01606   0.02666  -0.602   0.5469
## Lag4        -0.02779   0.02646  -1.050   0.2937
```

```

## Lag5      -0.01447   0.02638  -0.549   0.5833
## Volume    -0.02274   0.03690  -0.616   0.5377
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 1496.2 on 1088 degrees of freedom
## Residual deviance: 1486.4 on 1082 degrees of freedom
## AIC: 1500.4
##
## Number of Fisher Scoring iterations: 4

```

```
coef(glm.fit)
```

```

## (Intercept)      Lag1      Lag2      Lag3      Lag4      Lag5
##  0.26686414 -0.04126894  0.05844168 -0.01606114 -0.02779021 -0.01447206
##       Volume
## -0.02274153

```

Based on the p-value, only Lag2 is statistically significant.

(c)

```

glm.prob = predict(glm.fit, type = "response")
glm.prob[1:12] #Predicts P(Y=1/X), where 2 is the dummy variable for a rise in the stock market.

```

```

##      1      2      3      4      5      6      7
## 0.6086249 0.6010314 0.5875699 0.4816416 0.6169013 0.5684190 0.5786097
##      8      9     10     11     12
## 0.5151972 0.5715200 0.5554287 0.6092096 0.5370125

```

```
contrasts(Direction)
```

```

##      Up
## Down 0
## Up   1

```

```

glm.pred = rep("Down", 1089)
glm.pred[glm.prob>0.5] = "Up"
table(glm.pred, Direction)

```

```

##          Direction
## glm.pred Down Up
##      Down 54 48
##      Up   430 557

```

```
mean(glm.pred==Direction)
```

```
## [1] 0.5610652
```

The confusion matrix presents the wrongly classified data off the “left to right” diagonal matrix. In this scenario, the misclassified observations are 48 and 430, which sums up to 478 observations. However, the model accuracy (i.e., approx. 56.1%) does a slightly better job than a random guess (i.e., 50.0%).

(d)

```
train = (Year<2009)
Weekly.9.10 = Weekly[!train,]
dim(Weekly.9.10)

## [1] 104    9

Direction.9.10 = Direction[!train]

glm.fits = glm(Direction~Lag2, data = Weekly,
               family = binomial, subset = train)
glm.probs = predict(glm.fits, Weekly.9.10, type = "response")
glm preds = rep("Down", 104)
glm preds[glm.probs>.5] = "Up"
table(glm preds, Direction.9.10)
```

```
##          Direction.9.10
## glm preds Down Up
##      Down     9  5
##      Up      34 56
```

```
mean(glm preds == Direction.9.10) #model accuracy
```

```
## [1] 0.625
```

```
mean(glm preds != Direction.9.10) #Test error
```

```
## [1] 0.375
```

The overall model accuracy is 62.5%, and the test error is 37.5%. By utilizing the only significant predictor (i.e., Lag2) in the logistic regression model, model accuracy increased by approx. 6.5%.

The confusion matrix suggests that the accuracy of prediction of the upward and downward trend in the stock market are approx. 92% and 21%, respectively.

(e)

```

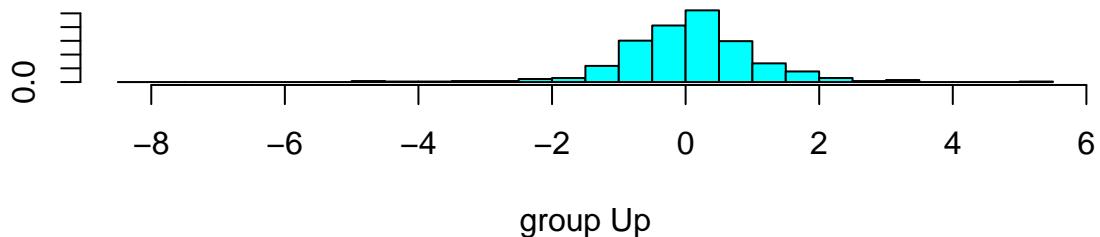
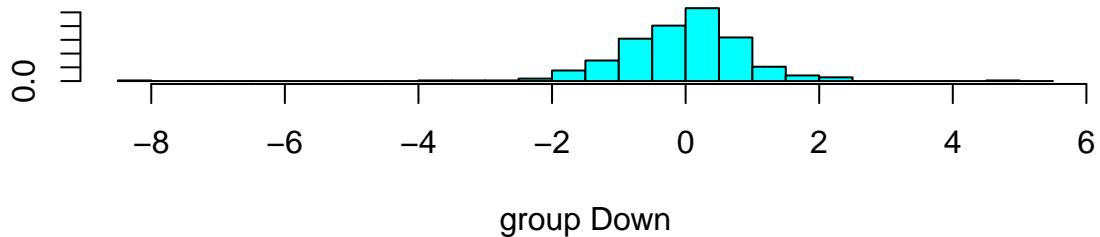
library(MASS)

lda.fit = lda(Direction~Lag2, data = Weekly, subset=train)
lda.fit

## Call:
## lda(Direction ~ Lag2, data = Weekly, subset = train)
##
## Prior probabilities of groups:
##       Down      Up
## 0.4477157 0.5522843
##
## Group means:
##           Lag2
## Down -0.03568254
## Up   0.26036581
##
## Coefficients of linear discriminants:
##          LD1
## Lag2 0.4414162

plot(lda.fit)

```



```
lda.pred = predict(lda.fit, Weekly.9.10)
names(lda.pred)
```

```
## [1] "class"      "posterior"   "x"
```

```
lda.class = lda.pred$class
table(lda.class, Direction.9.10)
```

```
##          Direction.9.10
## lda.class Down Up
##       Down    9  5
##       Up     34 56
```

```
mean(lda.class==Direction.9.10)
```

```
## [1] 0.625
```

Model accuracy is 62.5%, which is identical with the result obtained in the Logistic Regression that considered only Lag2 as the predictor.

(f)

```
qda.fit = qda(Direction~Lag2, data = Weekly, subset = train)
qda.fit
```

```
## Call:
## qda(Direction ~ Lag2, data = Weekly, subset = train)
##
## Prior probabilities of groups:
##       Down        Up
## 0.4477157 0.5522843
##
## Group means:
##           Lag2
## Down -0.03568254
## Up   0.26036581
```

```
qda.pred = predict(qda.fit, Weekly.9.10)
qda.class = qda.pred$class
table(qda.class,Direction.9.10)
```

```
##          Direction.9.10
## qda.class Down Up
##       Down    0  0
##       Up     43 61
```

```
mean(qda.class==Direction.9.10)
```

```
## [1] 0.5865385
```

The confusion matrix suggest that the model is 100% accurate in predicting market rise, but completely inaccurate (i.e., 0%) when predicting market fall.

Model accuracy is approx 58.7%.

(g)

```
library(class)
train.X= as.matrix(Lag2[train])
test.X= as.matrix(Lag2[!train])
train.Direction=Direction[train]
set.seed(7)
knn.pred=knn(train.X,test.X,train.Direction,k=1)
table(knn.pred,Direction.9.10)
```

```
##          Direction.9.10
## knn.pred Down Up
##       Down   21 30
##       Up     22 31
```

```
mean(knn.pred==Direction.9.10)
```

```
## [1] 0.5
```

The confusion matrix suggest that the model is approx 50.8% and 48.8% accurate in predicting the market rise and fall.

Model accuracy is approx. 50%.

(h)

With respect to overall model accuracy, Logistic regression and LDA provided the best results, followed by QDA and KNN-1.

h(i) - Freestyle

Logistic regression with Lag2 and Lag2~Lag1 synergy effect.

```
fit.glm1 = glm(Direction ~ Lag2+Lag2:Lag1, data = Weekly, family = binomial, subset = train)
glm.prob1 = predict(fit.glm1, Weekly.9.10, type = "response")
glm.pred1 = rep("Down", length(glm.prob1))
glm.pred1[glm.prob1>.5] = "Up"
table(glm.pred1, Direction.9.10)
```

```

##          Direction.9.10
## glm.pred1 Down Up
##      Down    3  3
##      Up     40 58

mean(glm.pred1==Direction.9.10)

```

[1] 0.5865385

LDA with Lag2 and Lag2~Lag1 synergy effect.

```

lda.fit1 = lda(Direction ~ Lag2+Lag2:Lag1, data = Weekly, subset = train)
lda.pred1 = predict(lda.fit1, Weekly.9.10)
table(lda.pred1$class, Direction.9.10)

```

```

##          Direction.9.10
##          Down Up
##      Down    3  3
##      Up     40 58

```

```
mean(lda.pred1$class == Direction.9.10)
```

[1] 0.5865385

Model accuracy is equivalent to the logistic regression model accuracy (i.e., 58.65%).

QDA with Lag2 and square of (abs(Lag2))

```

qda.fit1 = qda(Direction ~ Lag2 + (abs(Lag2))^2, data = Weekly, subset = train)
qda.pred1 = predict(qda.fit1, Weekly.9.10)
table(qda.pred1$class, Direction.9.10)

```

```

##          Direction.9.10
##          Down Up
##      Down   11 12
##      Up    32 49

```

```
mean(qda.pred1$class==Direction.9.10)
```

[1] 0.5769231

KNN (k = 5)

```

knn.pred1 = knn(train.X, test.X, train.Direction, k = 5)
table(knn.pred1, Direction.9.10)

```

```

##          Direction.9.10
## knn.pred1 Down Up
##      Down   15 20
##      Up    28 41

```

```
mean(knn.pred1==Direction.9.10)
```

```
## [1] 0.5384615
```

Model accuracy is 53.85%.

KNN (k = 10)

```
knn.pred2 = knn(train.X, test.X, train.Direction, k = 10)
table(knn.pred2, Direction.9.10)
```

```
##          Direction.9.10
## knn.pred2 Down Up
##      Down   18 20
##      Up     25 41
```

```
mean(knn.pred2==Direction.9.10)
```

```
## [1] 0.5673077
```

Logistic regression and LDA outperformed the other analysis methods, based on model accuracy.

Question 11

(a) - Attach mpg01

```
attach(Auto)

## The following object is masked from package:ggplot2:
##
##     mpg

mpg01 = rep(0, length(mpg))
mpg01[median(mpg)<mpg] = 1
Auto = data.frame(Auto, mpg01)
```

(b) - Explore the data

```
dim(Auto)
```

```
## [1] 392 10
```

```
str(Auto)
```

```
## 'data.frame': 392 obs. of 10 variables:
##   $ mpg      : num  18 15 18 16 17 15 14 14 14 15 ...
##   $ cylinders: num  8 8 8 8 8 8 8 8 8 ...
##   $ displacement: num  307 350 318 304 302 429 454 440 455 390 ...
##   $ horsepower: num  130 165 150 150 140 198 220 215 225 190 ...
##   $ weight    : num  3504 3693 3436 3433 3449 ...
##   $ acceleration: num  12 11.5 11 12 10.5 10 9 8.5 10 8.5 ...
##   $ year      : num  70 70 70 70 70 70 70 70 70 70 ...
##   $ origin    : num  1 1 1 1 1 1 1 1 1 ...
##   $ name      : Factor w/ 304 levels "amc ambassador brougham",...: 49 36 231 14 161 141 54 223 241 ...
##   $ mpg01     : num  0 0 0 0 0 0 0 0 0 0 ...
```

```
summary(Auto)
```

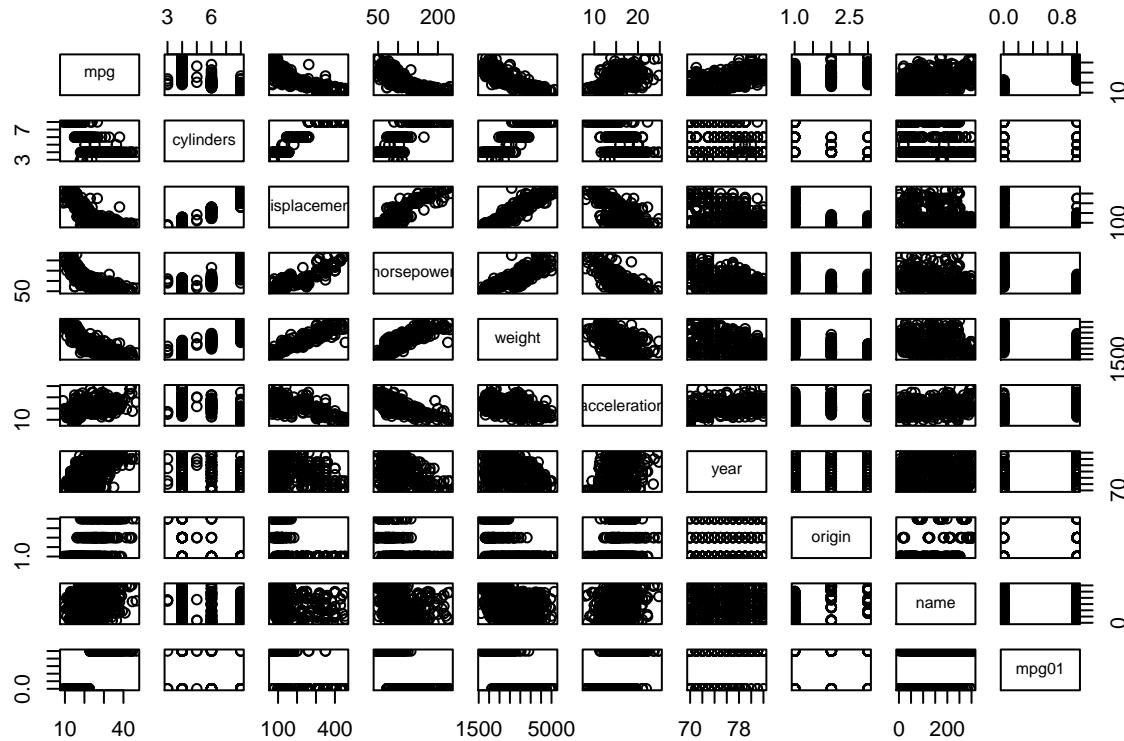
```
##          mpg           cylinders       displacement      horsepower
##  Min.   : 9.00   Min.   :3.000   Min.   :68.0   Min.   :46.0
##  1st Qu.:17.00  1st Qu.:4.000   1st Qu.:105.0  1st Qu.:75.0
##  Median :22.75  Median :4.000   Median :151.0  Median :93.5
##  Mean   :23.45  Mean   :5.472   Mean   :194.4  Mean   :104.5
##  3rd Qu.:29.00  3rd Qu.:8.000   3rd Qu.:275.8  3rd Qu.:126.0
##  Max.   :46.60  Max.   :8.000   Max.   :455.0  Max.   :230.0
##
##          weight        acceleration         year          origin
##  Min.   :1613   Min.   : 8.00   Min.   :70.00  Min.   :1.000
##  1st Qu.:2225  1st Qu.:13.78  1st Qu.:73.00  1st Qu.:1.000
```

```

## Median :2804   Median :15.50   Median :76.00   Median :1.000
## Mean    :2978   Mean    :15.54   Mean    :75.98   Mean    :1.577
## 3rd Qu.:3615   3rd Qu.:17.02   3rd Qu.:79.00   3rd Qu.:2.000
## Max.    :5140   Max.    :24.80   Max.    :82.00   Max.    :3.000
##
##          name      mpg01
## amc matador     : 5   Min.  :0.0
## ford pinto      : 5   1st Qu.:0.0
## toyota corolla   : 5   Median :0.5
## amc gremlin      : 4   Mean   :0.5
## amc hornet       : 4   3rd Qu.:1.0
## chevrolet chevette: 4   Max.   :1.0
## (Other)          :365

```

`pairs(Auto)`



`cor(Auto[, -9])`

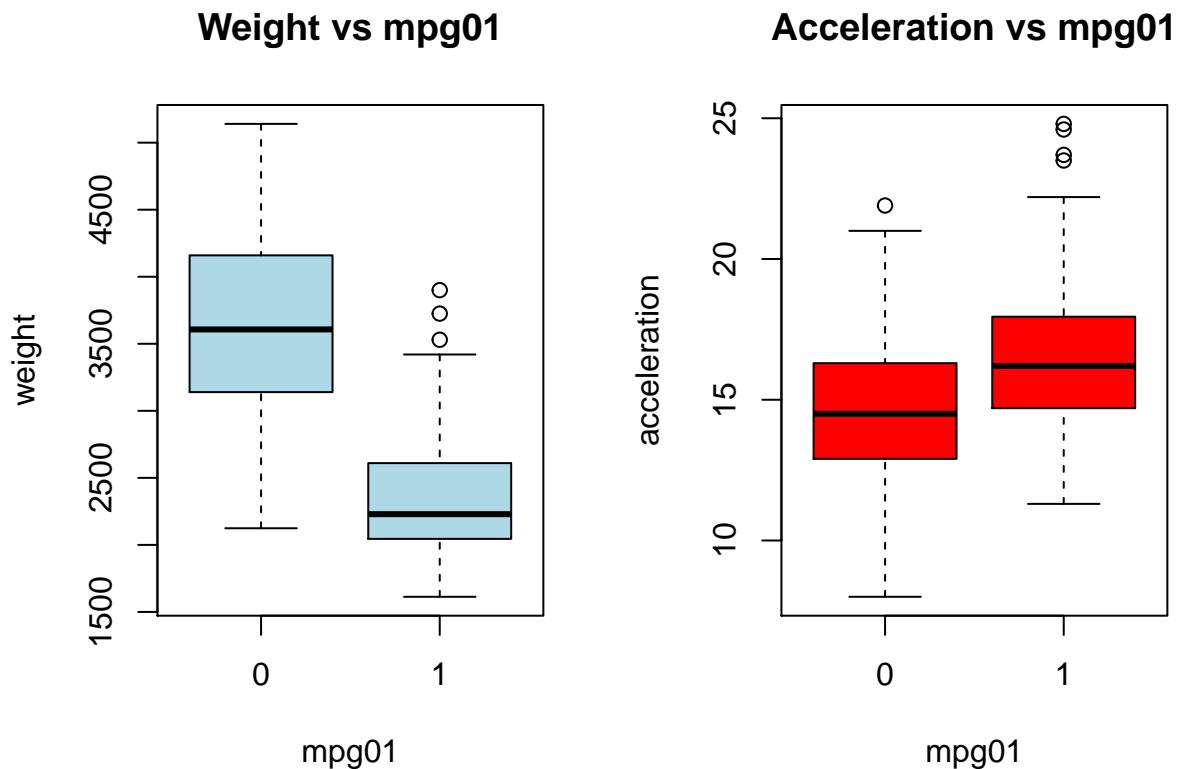
	mpg	cylinders	displacement	horsepower	weight
## mpg	1.0000000	-0.7776175	-0.8051269	-0.7784268	-0.8322442
## cylinders	-0.7776175	1.0000000	0.9508233	0.8429834	0.8975273
## displacement	-0.8051269	0.9508233	1.0000000	0.8972570	0.9329944
## horsepower	-0.7784268	0.8429834	0.8972570	1.0000000	0.8645377
## weight	-0.8322442	0.8975273	0.9329944	0.8645377	1.0000000
## acceleration	0.4233285	-0.5046834	-0.5438005	-0.6891955	-0.4168392

```

## year          0.5805410 -0.3456474   -0.3698552 -0.4163615 -0.3091199
## origin        0.5652088 -0.5689316   -0.6145351 -0.4551715 -0.5850054
## mpg01         0.8369392 -0.7591939   -0.7534766 -0.6670526 -0.7577566
## acceleration    0.4233285  0.5805410   0.5652088  0.8369392
## mpg            0.4233285  0.5805410   0.5652088  0.8369392
## cylinders      -0.5046834 -0.3456474   -0.5689316 -0.7591939
## displacement   -0.5438005 -0.3698552   -0.6145351 -0.7534766
## horsepower     -0.6891955 -0.4163615   -0.4551715 -0.6670526
## weight          -0.4168392 -0.3091199   -0.5850054 -0.7577566
## acceleration    1.0000000  0.2903161   0.2127458  0.3468215
## year           0.2903161  1.0000000   0.1815277  0.4299042
## origin          0.2127458  0.1815277   1.0000000  0.5136984
## mpg01           0.3468215  0.4299042   0.5136984  1.0000000

par(mfrow = c(1,2))
boxplot(weight~mpg01, data = Auto, main = "Weight vs mpg01", col = "light blue")
boxplot(acceleration~mpg01, data = Auto, main = "Acceleration vs mpg01", col = "red")

```

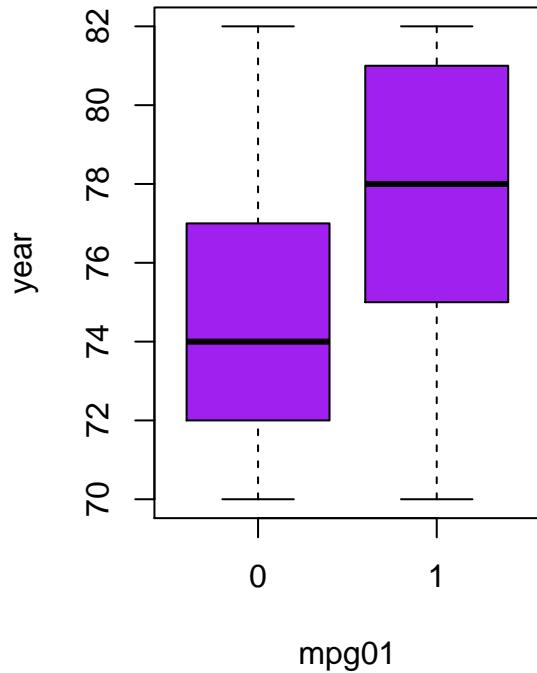


```

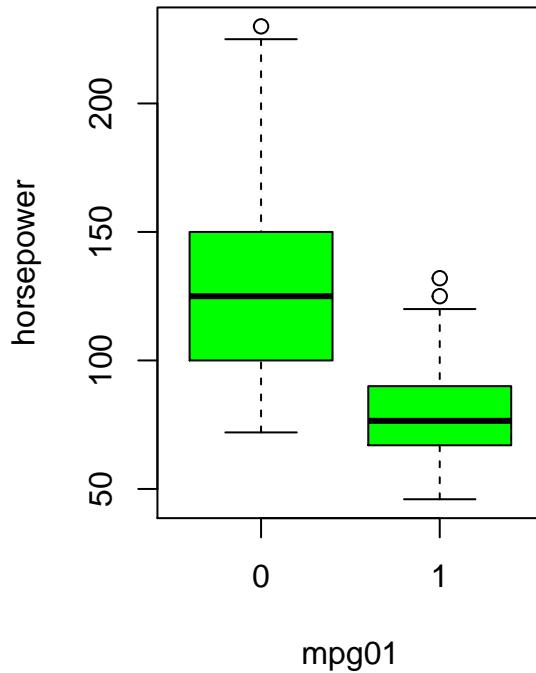
par(mfrow = c(1,2))
boxplot(year~mpg01, data = Auto, main = "Year vs mpg01", col = "purple")
boxplot(horsepower~mpg01, data = Auto, main = "Horsepower vs mpg01", col = "green")

```

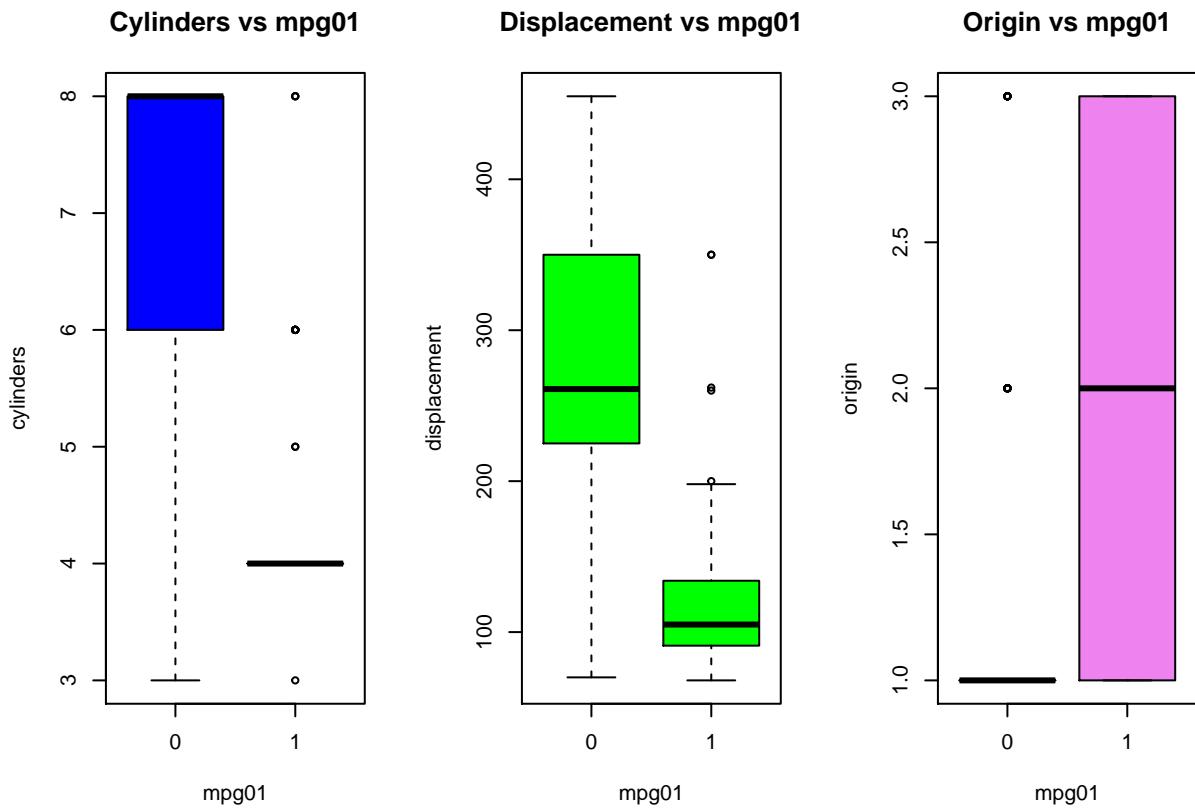
Year vs mpg01



Horsepower vs mpg01



```
par(mfrow = c(1,3))
boxplot(cylinders~mpg01, data = Auto, main = "Cylinders vs mpg01", col = "Blue")
boxplot(displacement~mpg01, data = Auto, main = "Displacement vs mpg01", col = "green")
boxplot(origin~mpg01, data = Auto, main = "Origin vs mpg01", col = "Violet")
```



From the boxplots, we can conclude that some relationships exist between “mpg01” and “weight,” “horsepower,” “year,” “acceleration,” and “displacement.”

(c) - split the data into training and testing sets

```
train = (year %% 2 == 0)
Auto.train = Auto[train, ]
Auto.test = Auto[!train, ]
mpg01.test <- mpg01[!train]
```

(d) - LDA

```
lda.fit2 = lda(mpg01 ~ weight+acceleration+horsepower+displacement+year,
               data = Auto, subset = train)
lda.fit2

## Call:
## lda(mpg01 ~ weight + acceleration + horsepower + displacement +
##       year, data = Auto, subset = train)
##
## Prior probabilities of groups:
##          0           1
```

```

## 0.4571429 0.5428571
##
## Group means:
##      weight acceleration horsepower displacement      year
## 0 3604.823      14.47500   133.14583     271.7396 74.10417
## 1 2314.763      16.62895    77.92105     111.6623 77.78947
##
## Coefficients of linear discriminants:
##                               LD1
## weight           -0.001583211
## acceleration   0.002468036
## horsepower      0.014846052
## displacement   -0.008311844
## year            0.105990658

lda.pred2 = predict(lda.fit2, Auto.test)
table(lda.pred2$class, mpg01.test)

```

```

##      mpg01.test
##      0 1
## 0 84 6
## 1 16 76

mean(lda.pred2$class != mpg01.test)

```

```
## [1] 0.1208791
```

Test error is approx.12.1%.

(e) - QDA

```

qda.fit2 = qda(mpg01 ~ weight+acceleration+horsepower+displacement+year,
                data = Auto, subset = train)
qda.fit2

```

```

## Call:
## qda(mpg01 ~ weight + acceleration + horsepower + displacement +
##       year, data = Auto, subset = train)
##
## Prior probabilities of groups:
##          0          1
## 0.4571429 0.5428571
##
## Group means:
##      weight acceleration horsepower displacement      year
## 0 3604.823      14.47500   133.14583     271.7396 74.10417
## 1 2314.763      16.62895    77.92105     111.6623 77.78947

```

```
qda.pred2 = predict(qda.fit2,Auto.test)
table(qda.pred2$class,mpg01.test)
```

```
##      mpg01.test
##      0   1
##  0 89  8
##  1 11 74
```

```
mean(qda.pred2$class!=mpg01.test)
```

```
## [1] 0.1043956
```

Test error is approx.10.4%.

(f) - Logistic Regression

```
glm.fit2 = glm(mpg01 ~ weight+acceleration+horsepower+displacement+year,
               data = Auto, family = binomial, subset = train)
summary(glm.fit2)
```

```
##
## Call:
## glm(formula = mpg01 ~ weight + acceleration + horsepower + displacement +
##       year, family = binomial, data = Auto, subset = train)
##
## Deviance Residuals:
##    Min      1Q  Median      3Q     Max
## -2.4415 -0.0059  0.0142  0.1366  1.7122
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept) -11.044917  8.453730 -1.307  0.1914
## weight       -0.004446  0.001850 -2.404  0.0162 *
## acceleration -0.180881  0.251599 -0.719  0.4722
## horsepower   -0.069563  0.037632 -1.849  0.0645 .
## displacement -0.021693  0.011561 -1.876  0.0606 .
## year          0.478827  0.118756  4.032 5.53e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 289.577  on 209  degrees of freedom
## Residual deviance: 57.762  on 204  degrees of freedom
## AIC: 69.762
##
## Number of Fisher Scoring iterations: 8
```

```

glm.prob2 = predict(glm.fit2, Auto.test, type = "response")
glm.pred2 = rep(0, length(glm.prob2))
glm.pred2[glm.prob2 > 0.5] = 1
table(glm.pred2, mpg01.test)

##           mpg01.test
## glm.pred2  0   1
##             0 88  8
##             1 12 74

mean(glm.pred2 != mpg01.test)

```

[1] 0.1098901

Test error is approx. 11.0%.

(g) - KNN (K = 1)

```

train.X1 = cbind(weight,acceleration,horsepower,displacement,year)[train,]
test.X1 = cbind(weight,acceleration,horsepower,displacement,year)[!train,]
train.mpg01 = mpg01[train]
set.seed(7)
knn.pred3 = knn(train.X1, test.X1, train.mpg01, k = 1)
table(knn.pred3, mpg01.test)

```

```

##           mpg01.test
## knn.pred3  0   1
##             0 83 11
##             1 17 71

mean(knn.pred3 != mpg01.test)

```

[1] 0.1538462

Test error is approx. 15.4%.

K = 10

```

train.X1 = cbind(weight,acceleration,horsepower,displacement,year)[train,]
test.X1 = cbind(weight,acceleration,horsepower,displacement,year)[!train,]
train.mpg01 = mpg01[train]
set.seed(7)
knn.pred4 = knn(train.X1, test.X1, train.mpg01, k = 10)
table(knn.pred4, mpg01.test)

```

```

##           mpg01.test
## knn.pred4  0   1
##             0 76  7
##             1 24 75

```

```
mean(knn.pred4 != mpg01.test)
```

```
## [1] 0.1703297
```

Test error is approx. 15.9%.

K = 100

```
train.X1 = cbind(weight,acceleration,horsepower,displacement,year)[train,]
test.X1 = cbind(weight,acceleration,horsepower,displacement,year)[!train,]
train.mpg01 = mpg01[train]
set.seed(7)
knn.pred5 = knn(train.X1, test.X1, train.mpg01, k = 100)
table(knn.pred5, mpg01.test)
```

```
##          mpg01.test
## knn.pred5 0 1
##                 0 81 7
##                 1 19 75
```

```
mean(knn.pred5 != mpg01.test)
```

```
## [1] 0.1428571
```

Test error is approx. 14.3%. A K value of 100 seems to yield the lowest test error.